

# capstoneProject

December 1, 2020

## 1 Capstone Project

### 1.1 Image classifier for the SVHN dataset

#### 1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### 1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

#### 1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
[17]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
from scipy.io import loadmat
from random import randint
from PIL import Image
```

For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. “Reading Digits in Natural Images with Unsupervised Feature Learning”. NIPS Workshop on Deep Learning

and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

```
[18]: # Run this cell to load the dataset

train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

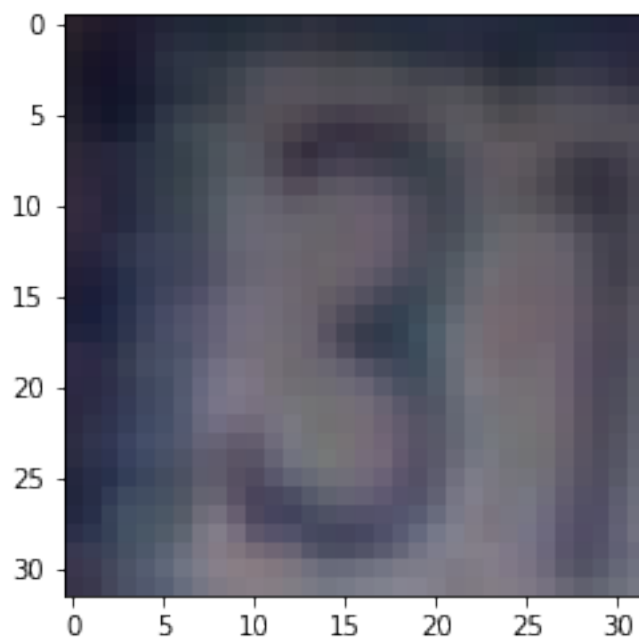
Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

## 1.2 1. Inspect and preprocess the dataset

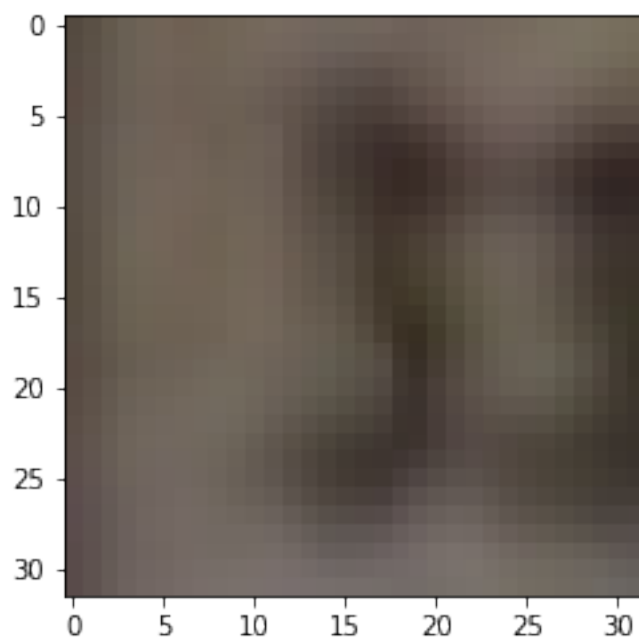
- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
[19]: train_img_arr = train['X']
train_img_labels = train['y']
test_img_arr = test['X']
test_img_labels = test['y']
```

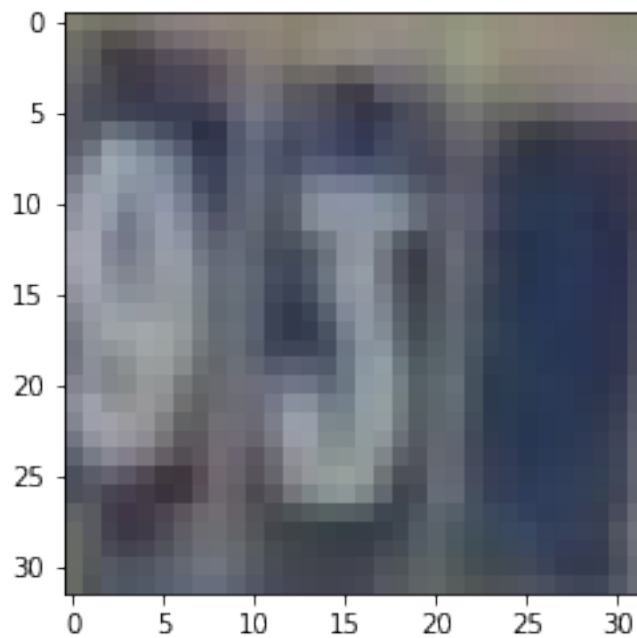
```
[20]: lenOfTrain = len(train_img_arr)
for x in range(10):
    randomInt = randint(0, lenOfTrain-1)
    plt.imshow(train_img_arr[:, :, :, randomInt])
    plt.show()
    print("img sample label: ", train_img_labels[randomInt])
```



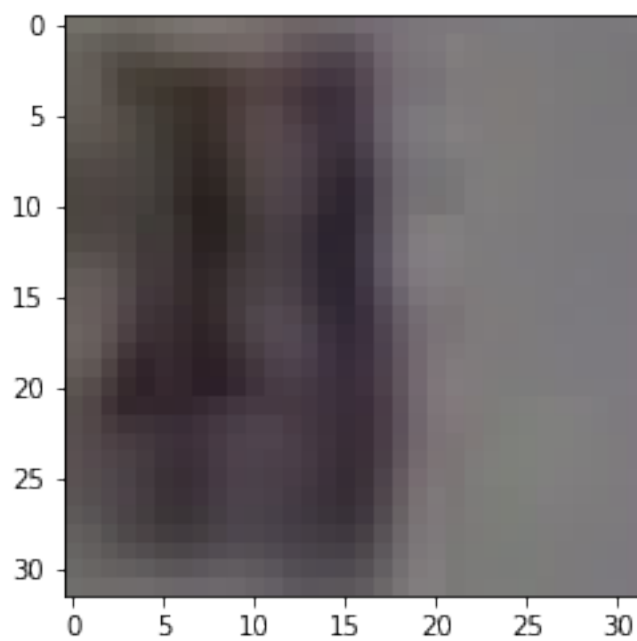
img sample label: [3]



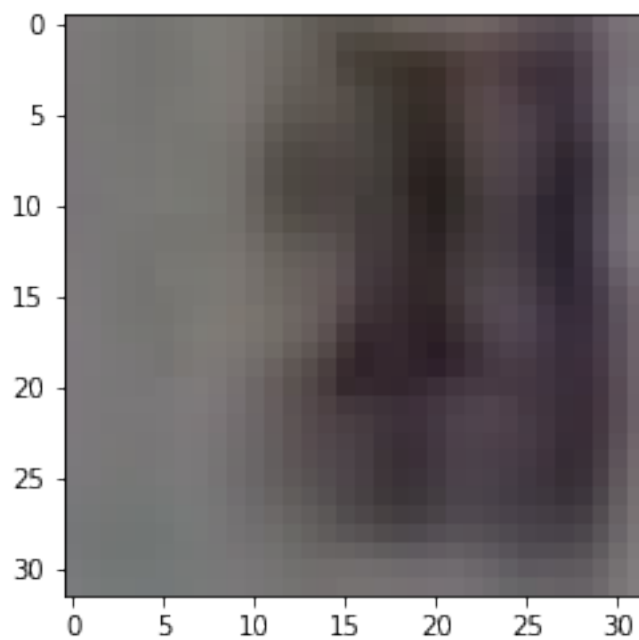
img sample label: [3]



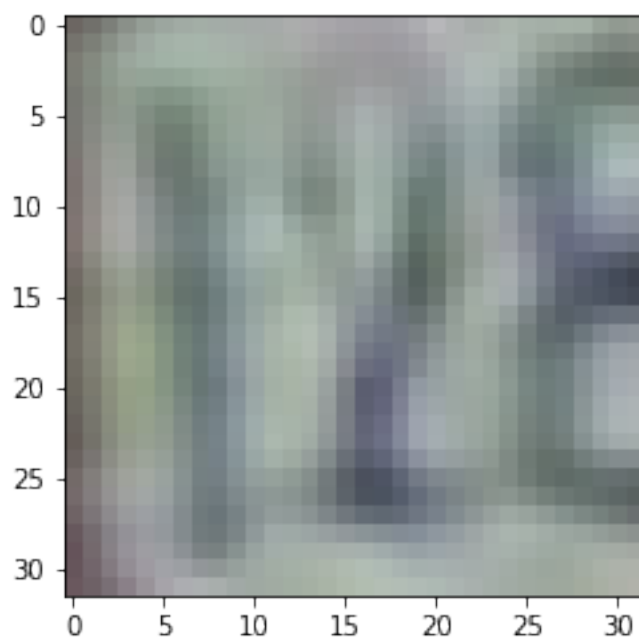
img sample label: [3]



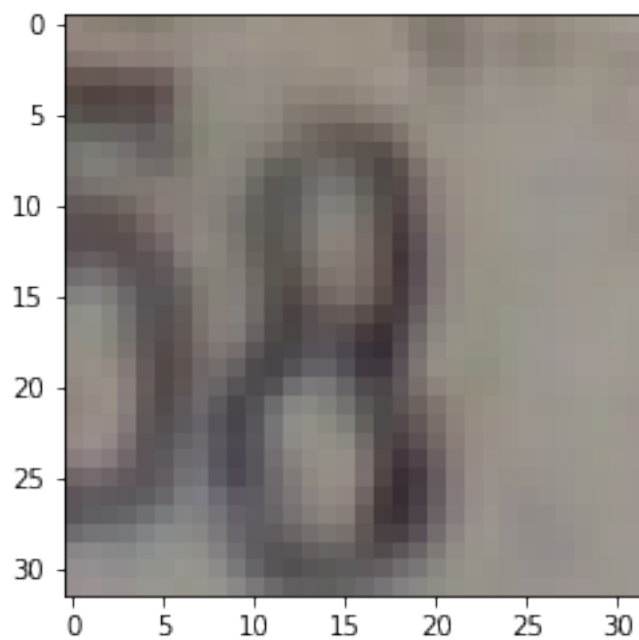
img sample label: [3]



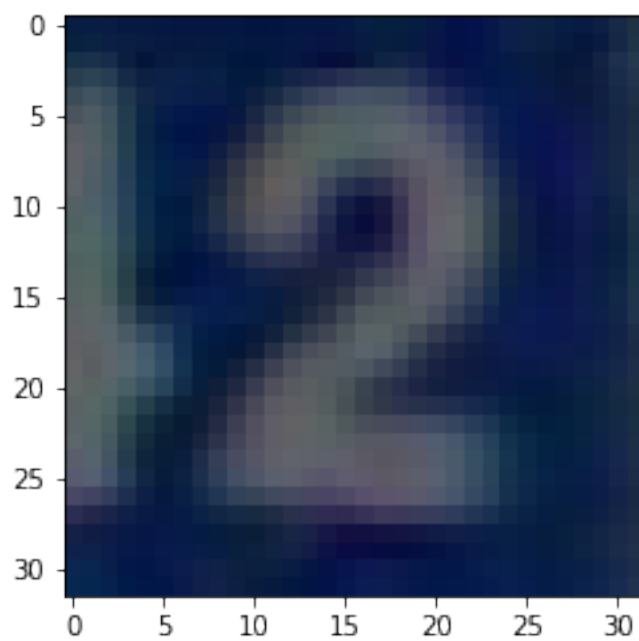
img sample label: [2]



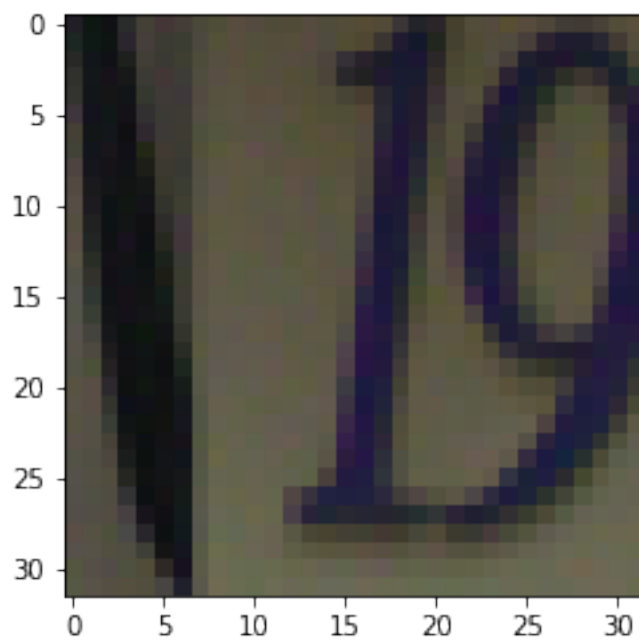
img sample label: [2]



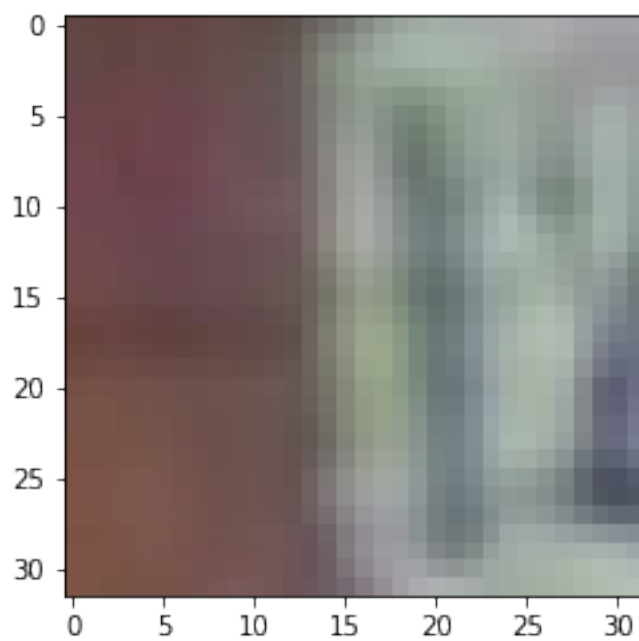
img sample label: [8]



img sample label: [2]



img sample label: [1]



img sample label: [1]

```
[21]: # Converts input RGB array to grayscale equivalent

def convert_arr_to_gray(input_array):
    # get shape of train img arr
    original_shape = np.shape(input_array)
    # create gray scale equivalent array with one channel.
    # change the format to have the batch the first parm
    img_arr_gray = np.zeros((original_shape[3], original_shape[0],
→original_shape[1], 1))
    for x in range(np.shape(input_array)[0]):
        for y in range(np.shape(input_array)[1]):
            # convert rgb to grayscale
            gray_arr = ((input_array[x][y][0] + input_array[x][y][1] +
→input_array[x][y][2]) / 3.0) / 255.
            # set it in output arr
            img_arr_gray[:,x,y,0] = gray_arr
    print(img_arr_gray)
    return img_arr_gray

train_img_arr_gray = convert_arr_to_gray(train_img_arr)
print(f"shape: {np.shape(train_img_arr_gray)}")
test_img_arr_gray = convert_arr_to_gray(test_img_arr)
```

```
[[[0.13202614]
    [0.0745098 ]
    [0.06666667]
    ...
    [0.25228758]
    [0.28235294]
    [0.29673203]]

  [[0.13333333]
    [0.07973856]
    [0.06666667]
    ...
    [0.16862745]
    [0.23137255]
    [0.28888889]]

  [[0.15555556]
    [0.08496732]
    [0.07320261]
    ...
    [0.12679739]
    [0.15816993]
    [0.25359477]]

  ...
```



[0.31633987]  
[0.31372549]  
[0.29542484]  
...  
[0.05098039]  
[0.04444444]  
[0.03921569]]

[0.30588235]  
[0.30980392]  
[0.30065359]  
...  
[0.05359477]  
[0.05228758]  
[0.05228758]]

[0.31503268]  
[0.31764706]  
[0.30196078]  
...  
[0.06013072]  
[0.05751634]  
[0.04183007]]]

[[0.28627451]  
[0.29411765]  
[0.27581699]  
...  
[0.30980392]  
[0.30326797]  
[0.30326797]]

[0.29150327]  
[0.2875817 ]  
[0.26405229]  
...  
[0.30718954]  
[0.30980392]  
[0.30849673]]

[0.28627451]  
[0.27843137]  
[0.21568627]  
...  
[0.31111111]  
[0.30980392]

```

[0.30980392]]

...

[[0.01830065]
 [0.01568627]
 [0.00653595]
 ...
 [0.0496732 ]
 [0.04575163]
 [0.03006536]]

[[0.03921569]
 [0.04313725]
 [0.04836601]
 ...
 [0.04836601]
 [0.04836601]
 [0.0379085 ]]

[[0.04575163]
 [0.05098039]
 [0.05882353]
 ...
 [0.0627451 ]
 [0.05228758]
 [0.03267974]]]

[[[0.23921569]
 [0.23921569]
 [0.25882353]
 ...
 [0.23006536]
 [0.24575163]
 [0.2875817 ]]]

[[0.24052288]
 [0.23137255]
 [0.25098039]
 ...
 [0.19869281]
 [0.26405229]
 [0.03529412]]

[[0.24052288]
 [0.23137255]
 [0.23529412]

```

```

...
[0.23529412]
[0.03137255]
[0.18039216]]

...

[[0.30196078]
 [0.30588235]
 [0.27973856]

...
 [0.22875817]
 [0.23660131]
 [0.24575163]]

[[0.30849673]
 [0.29803922]
 [0.29803922]

...
 [0.22875817]
 [0.22745098]
 [0.23267974]]

[[0.29542484]
 [0.27712418]
 [0.29673203]

...
 [0.24313725]
 [0.23660131]
 [0.23921569]]]

...

[[[0.01960784]
  [0.03267974]
  [0.10980392]

...
  [0.12026144]
  [0.08888889]
  [0.08104575]]

[[0.02222222]
 [0.03006536]
 [0.10065359]

...
 [0.04313725]

```

```

[0.02745098]
[0.05228758]]

[[0.05098039]
 [0.0379085 ]
 [0.08888889]
 ...
 [0.32418301]
 [0.00915033]
 [0.06013072]]

...

[[0.06143791]
 [0.05228758]
 [0.10326797]
 ...
 [0.14248366]
 [0.1254902 ]
 [0.11372549]]

[[0.04183007]
 [0.04313725]
 [0.10718954]
 ...
 [0.18300654]
 [0.17385621]
 [0.15424837]]

[[0.01699346]
 [0.00784314]
 [0.08627451]
 ...
 [0.17777778]
 [0.19346405]
 [0.18431373]]]

[[[0.0745098 ]
 [0.13333333]
 [0.19084967]
 ...
 [0.2248366 ]
 [0.22222222]
 [0.21960784]]

[[0.05098039]
 [0.11503268]

```

```

[0.18431373]
...
[0.21699346]
[0.2130719 ]
[0.21045752]]

[[0.00522876]
 [0.07712418]
 [0.16862745]
 ...
 [0.22875817]
 [0.21699346]
 [0.20653595]]

...

[[0.06013072]
 [0.12156863]
 [0.17254902]
 ...
 [0.11372549]
 [0.08627451]
 [0.08627451]]

[[0.05620915]
 [0.10065359]
 [0.14901961]
 ...
 [0.18039216]
 [0.15163399]
 [0.13202614]]

[[0.07320261]
 [0.10196078]
 [0.1372549 ]
 ...
 [0.22222222]
 [0.19738562]
 [0.16732026]]]

[[[0.17385621]
   [0.19346405]
   [0.21176471]
   ...
   [0.10196078]
   [0.06666667]
   [0.07843137]]]

```

```

[[0.12941176]
 [0.15294118]
 [0.19084967]
 ...
 [0.10588235]
 [0.07843137]
 [0.07843137]]

[[0.10588235]
 [0.12156863]
 [0.15816993]
 ...
 [0.11372549]
 [0.09281046]
 [0.07843137]]

...

[[0.22875817]
 [0.21830065]
 [0.20915033]
 ...
 [0.09411765]
 [0.04183007]
 [0.06013072]]

[[0.22614379]
 [0.21699346]
 [0.20915033]
 ...
 [0.08235294]
 [0.03006536]
 [0.04444444]]

[[0.22352941]
 [0.22352941]
 [0.21960784]
 ...
 [0.06405229]
 [0.01568627]
 [0.02352941]]]]
shape: (73257, 32, 32, 1)
[[[[0.2627451 ]
 [0.26666667]
 [0.26797386]
 ...
 [0.26666667]

```

```

[0.27058824]
[0.25228758]]

[[0.26666667]
 [0.26666667]
 [0.26797386]
 ...
 [0.26666667]
 [0.27058824]
 [0.25228758]]

[[0.26797386]
 [0.26405229]
 [0.27189542]
 ...
 [0.27058824]
 [0.27058824]
 [0.25490196]]

...

[[0.3254902 ]
 [0.32941176]
 [0.32156863]
 ...
 [0.30196078]
 [0.30326797]
 [0.29281046]]

[[0.31764706]
 [0.32156863]
 [0.31372549]
 ...
 [0.30326797]
 [0.30457516]
 [0.29281046]]

[[0.30980392]
 [0.30980392]
 [0.30196078]
 ...
 [0.30065359]
 [0.30065359]
 [0.29542484]]]

[[[0.21960784]
 [0.21699346]

```

```

[0.2130719 ]
...
[0.19738562]
[0.20522876]
[0.20784314]]

[[0.24575163]
 [0.24052288]
 [0.2379085 ]
 ...
 [0.21699346]
 [0.21437908]
 [0.21437908]]

[[0.27581699]
 [0.27058824]
 [0.26928105]
 ...
 [0.24705882]
 [0.23529412]
 [0.22875817]]

...

[[0.29542484]
 [0.27581699]
 [0.25359477]
 ...
 [0.07843137]
 [0.08888889]
 [0.12156863]]

[[0.26013072]
 [0.22875817]
 [0.19477124]
 ...
 [0.08366013]
 [0.1124183 ]
 [0.15424837]]

[[0.23267974]
 [0.19738562]
 [0.15686275]
 ...
 [0.11503268]
 [0.15294118]
 [0.2      ]]]

```



[[0.29150327]  
[0.29673203]  
[0.30849673]  
...  
[0.3254902 ]  
[0.32156863]  
[0.31372549]]

[[0.29150327]  
[0.29673203]  
[0.30849673]  
...  
[0.33202614]  
[0.32418301]  
[0.31633987]]

[[0.26797386]  
[0.2745098 ]  
[0.28627451]  
...  
[0.32679739]  
[0.31372549]  
[0.30326797]]

...

[[0.25620915]  
[0.25228758]  
[0.24575163]  
...  
[0.21568627]  
[0.21176471]  
[0.21045752]]

[[0.30718954]  
[0.30588235]  
[0.30588235]  
...  
[0.27189542]  
[0.27189542]  
[0.27189542]]

[[0.01699346]  
[0.01830065]  
[0.01960784]  
...  
[0.32026144]

```
[0.32156863]
[0.32287582]]]
```

...

```
[[[0.17385621]
   [0.17777778]
   [0.18169935]
   ...
   [0.17647059]
   [0.17777778]
   [0.18039216]]]
```

```
[[0.2      ]
 [0.20392157]
 [0.20784314]
 ...
 [0.18823529]
 [0.19084967]
 [0.19346405]]]
```

```
[[0.23267974]
 [0.23660131]
 [0.23921569]
 ...
 [0.20784314]
 [0.21045752]
 [0.21437908]]]
```

...

```
[[0.06405229]
 [0.06013072]
 [0.05620915]
 ...
 [0.26013072]
 [0.26797386]
 [0.2745098  ]]
```

```
[[0.15686275]
 [0.15686275]
 [0.15947712]
 ...
 [0.26405229]
 [0.27058824]
 [0.2745098  ]]
```

[0.21699346]  
[0.22091503]  
[0.22745098]  
...  
[0.26797386]  
[0.27189542]  
[0.2745098 ]]]

[[0.27189542]  
[0.27581699]  
[0.27712418]  
...  
[0.2627451 ]  
[0.26666667]  
[0.27058824]]

[0.27320261]  
[0.27712418]  
[0.27843137]  
...  
[0.26143791]  
[0.26666667]  
[0.27058824]]

[0.27320261]  
[0.2745098 ]  
[0.27581699]  
...  
[0.27058824]  
[0.28104575]  
[0.28496732]]

...

[0.28104575]  
[0.28627451]  
[0.29542484]  
...  
[0.14117647]  
[0.09542484]  
[0.05490196]]

[0.27581699]  
[0.28366013]  
[0.29281046]  
...

```

[0.09150327]
[0.05751634]
[0.02614379]]

[[0.27189542]
 [0.28235294]
 [0.29150327]
 ...
 [0.05620915]
 [0.02745098]
 [0.00130719]]]

[[[0.30849673]
   [0.30457516]
   [0.29934641]
   ...
   [0.2745098 ]
   [0.2745098 ]
   [0.27581699]]

 [0.00784314]
 [0.33202614]
 [0.31895425]
 ...
 [0.2745098 ]
 [0.26666667]
 [0.26535948]]

 [0.04444444]
 [0.03006536]
 [0.00653595]
 ...
 [0.2745098 ]
 [0.26405229]
 [0.25882353]]

...

[[0.05359477]
 [0.05620915]
 [0.05359477]
 ...
 [0.28888889]
 [0.28888889]
 [0.28888889]]

[[0.00261438]

```

```

[0.00784314]
[0.00915033]
...
[0.2875817 ]
[0.29150327]
[0.29150327]]

[[0.30065359]
 [0.30457516]
 [0.30849673]
 ...
 [0.2875817 ]
 [0.29150327]
 [0.29150327]]]]

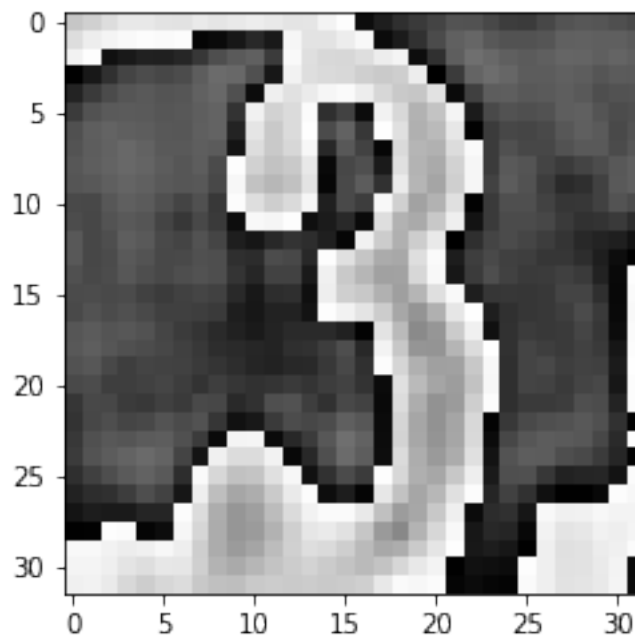
```

```

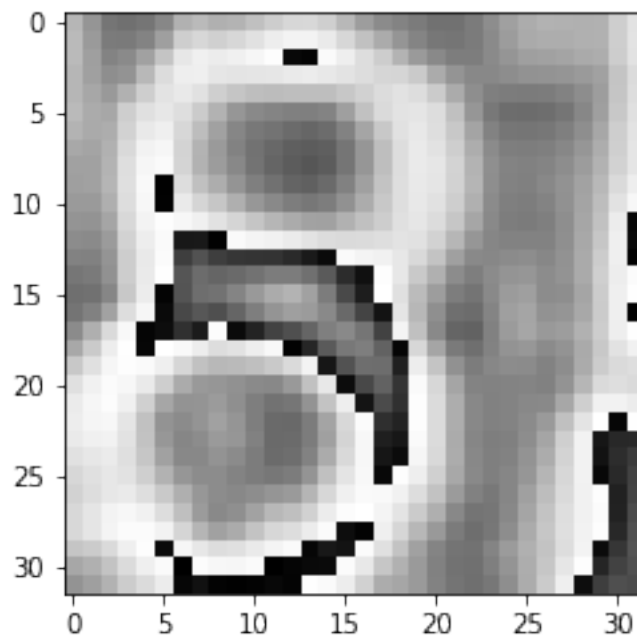
[22]: len_of_train_gray = len(train_img_arr_gray)
      print(np.shape(train_img_arr_gray))
      for x in range(10):
          randomInt = randint(0, len_of_train_gray-1)
          randomImage = train_img_arr_gray[randomInt,:,:,:]
          plt.imshow(randomImage[:,:,:0], cmap='Greys')
          plt.show()
          print("img sample label: ", train_img_labels[randomInt])

```

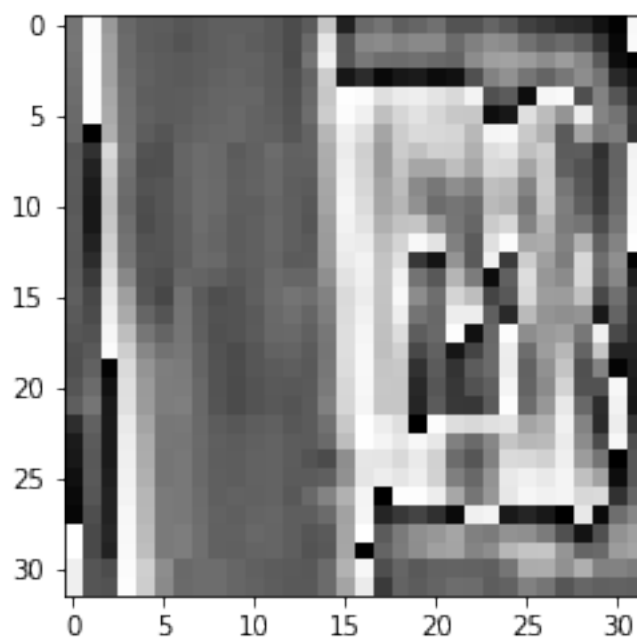
(73257, 32, 32, 1)



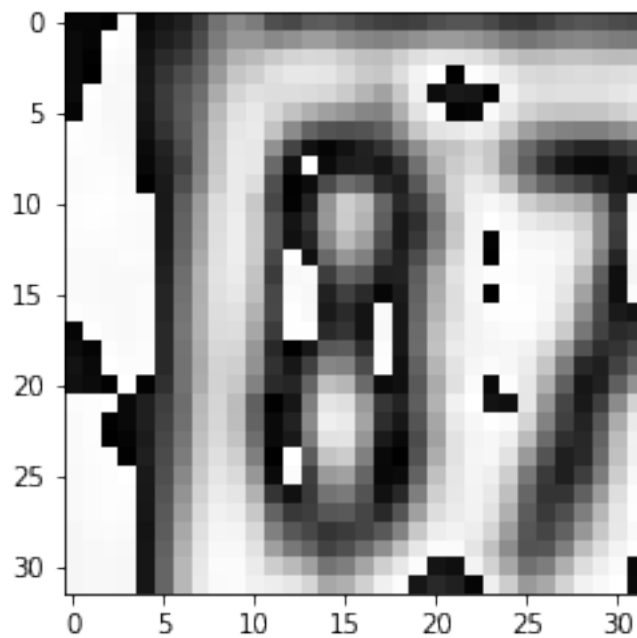
img sample label: [3]



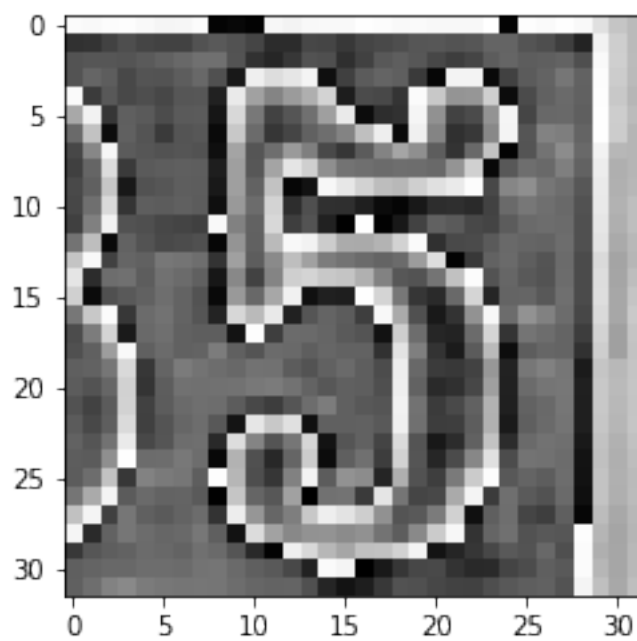
img sample label: [8]



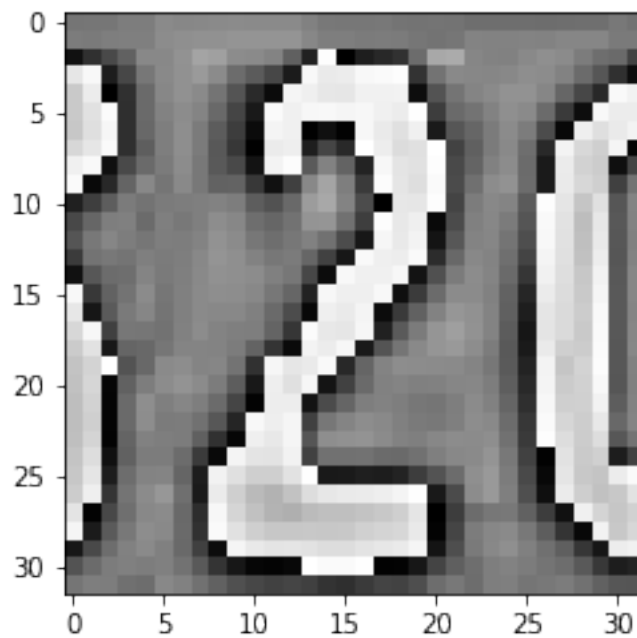
img sample label: [1]



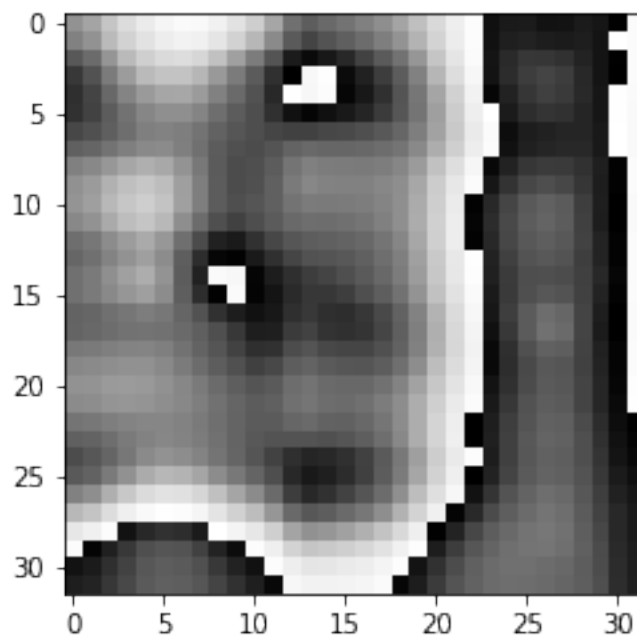
img sample label: [8]



img sample label: [5]

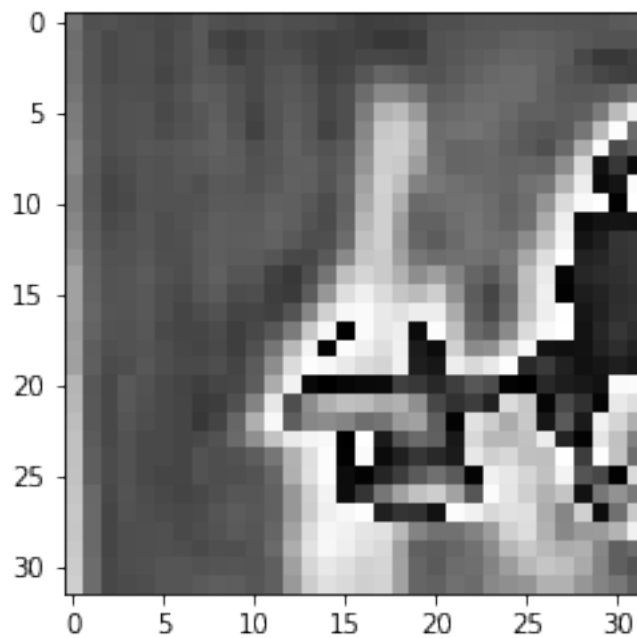


img sample label: [2]

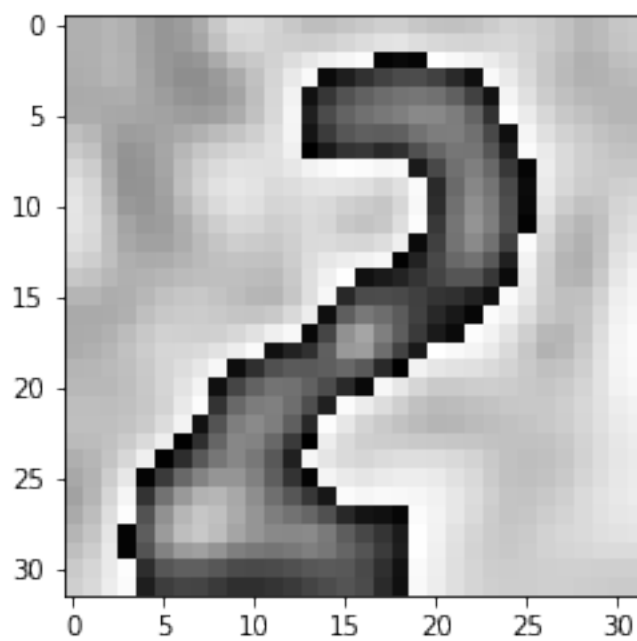


img sample label: [6]

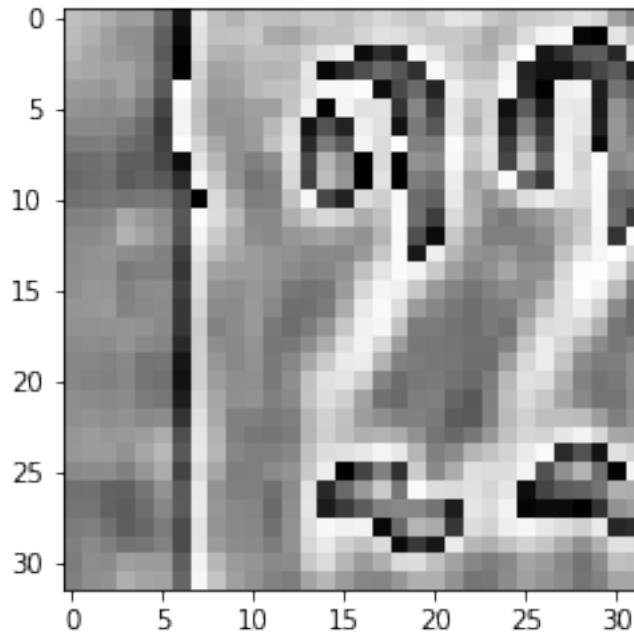




img sample label: [4]



img sample label: [2]



img sample label: [2]

[ ]:

[ ]:

### 1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[23]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense, Flatten, Softmax
```

```
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
```

```
[34]: # Build the Sequential feedforward neural network model
def get_model(input_shape):
    model = Sequential([
        Flatten(input_shape=input_shape),
        Dense(128, activation='relu', name='layer_1'),
        Dense(64, activation='relu', name='layer_2'),
        Dense(32, activation='relu', name='layer_3'),
        Dense(11, activation='softmax', name='output_layer')
    ])
    return model

input_shape = (np.shape(train_img_arr_gray)[1], np.
    ↳shape(train_img_arr_gray)[2], np.shape(train_img_arr_gray)[3])

model = get_model(input_shape)
```

```
[35]: # Print the model summary
```

```
model.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
flatten_3 (Flatten)	(None, 1024)	0
layer_1 (Dense)	(None, 128)	131200
layer_2 (Dense)	(None, 64)	8256
layer_3 (Dense)	(None, 32)	2080
output_layer (Dense)	(None, 11)	363

Total params: 141,899  
Trainable params: 141,899  
Non-trainable params: 0

```
[37]: #compile the model
acc = tf.keras.metrics.SparseCategoricalAccuracy()
```

```
model.compile(optimizer="Adam", loss='sparse_categorical_crossentropy',  
↳metrics=[acc])
```

```
[38]: def get_checkpoint_best_only(checkpoint_best_path):  
    """  
    Create best checkpoint callback  
    """  
    checkpoint_best = ModelCheckpoint(filepath=checkpoint_best_path,  
                                     save_weights_only=True,  
                                     save_freq='epoch',  
                                     monitor='sparse_categorical_accuracy',  
                                     save_best_only=True,  
                                     verbose=1)  
    return checkpoint_best  
  
def get_early_stopping():  
    """  
    Creates EarlyStopping callback that stops training when  
    the validation (testing) accuracy has not improved in the last 3 epochs.  
    """  
    return EarlyStopping(patience=2, monitor='sparse_categorical_accuracy')  
  
checkpoint_best_only_mlp = get_checkpoint_best_only('./  
↳checkpoints_mlp_best_only/checkpoint')  
early_stopping_mlp = get_early_stopping()  
callbacks_mlp = [checkpoint_best_only_mlp, early_stopping_mlp]
```

```
[39]: # Fit the model  
history = model.fit(train_img_arr_gray,  
                    train_img_labels,  
                    validation_split=0.15,  
                    epochs=30,  
                    batch_size=32,  
                    verbose=2,  
                    callbacks=callbacks_mlp)
```

Epoch 1/30

Epoch 00001: sparse\_categorical\_accuracy improved from -inf to 0.18673, saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 2.2466 - sparse\_categorical\_accuracy: 0.1867 - val\_loss: 2.2443 - val\_sparse\_categorical\_accuracy: 0.1881

Epoch 2/30

Epoch 00002: sparse\_categorical\_accuracy improved from 0.18673 to 0.20515, saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 2.1950 - sparse\_categorical\_accuracy: 0.2051 - val\_loss: 2.1313 - val\_sparse\_categorical\_accuracy: 0.2342

Epoch 3/30

Epoch 00003: sparse\_categorical\_accuracy improved from 0.20515 to 0.26909,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 2.0418 - sparse\_categorical\_accuracy: 0.2691 - val\_loss:  
1.9951 - val\_sparse\_categorical\_accuracy: 0.3028

Epoch 4/30

Epoch 00004: sparse\_categorical\_accuracy improved from 0.26909 to 0.32707,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.9298 - sparse\_categorical\_accuracy: 0.3271 - val\_loss:  
1.9019 - val\_sparse\_categorical\_accuracy: 0.3466

Epoch 5/30

Epoch 00005: sparse\_categorical\_accuracy improved from 0.32707 to 0.35818,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.8690 - sparse\_categorical\_accuracy: 0.3582 - val\_loss:  
1.8827 - val\_sparse\_categorical\_accuracy: 0.3525

Epoch 6/30

Epoch 00006: sparse\_categorical\_accuracy improved from 0.35818 to 0.37401,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.8294 - sparse\_categorical\_accuracy: 0.3740 - val\_loss:  
1.8472 - val\_sparse\_categorical\_accuracy: 0.3671

Epoch 7/30

Epoch 00007: sparse\_categorical\_accuracy improved from 0.37401 to 0.38933,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.7957 - sparse\_categorical\_accuracy: 0.3893 - val\_loss:  
1.8170 - val\_sparse\_categorical\_accuracy: 0.3803

Epoch 8/30

Epoch 00008: sparse\_categorical\_accuracy improved from 0.38933 to 0.40014,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.7663 - sparse\_categorical\_accuracy: 0.4001 - val\_loss:  
1.7945 - val\_sparse\_categorical\_accuracy: 0.3916

Epoch 9/30

Epoch 00009: sparse\_categorical\_accuracy improved from 0.40014 to 0.41318,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.7423 - sparse\_categorical\_accuracy: 0.4132 - val\_loss:  
1.8091 - val\_sparse\_categorical\_accuracy: 0.3883

Epoch 10/30

Epoch 00010: sparse\_categorical\_accuracy improved from 0.41318 to 0.42097,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.7196 - sparse\_categorical\_accuracy: 0.4210 - val\_loss:  
1.7697 - val\_sparse\_categorical\_accuracy: 0.3982

Epoch 11/30

Epoch 00011: sparse\_categorical\_accuracy improved from 0.42097 to 0.42943,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6983 - sparse\_categorical\_accuracy: 0.4294 - val\_loss:  
1.7699 - val\_sparse\_categorical\_accuracy: 0.4025

Epoch 12/30

Epoch 00012: sparse\_categorical\_accuracy improved from 0.42943 to 0.43562,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6802 - sparse\_categorical\_accuracy: 0.4356 - val\_loss:  
1.7510 - val\_sparse\_categorical\_accuracy: 0.4112

Epoch 13/30

Epoch 00013: sparse\_categorical\_accuracy improved from 0.43562 to 0.44498,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6633 - sparse\_categorical\_accuracy: 0.4450 - val\_loss:  
1.7514 - val\_sparse\_categorical\_accuracy: 0.4128

Epoch 14/30

Epoch 00014: sparse\_categorical\_accuracy improved from 0.44498 to 0.44747,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6509 - sparse\_categorical\_accuracy: 0.4475 - val\_loss:  
1.7235 - val\_sparse\_categorical\_accuracy: 0.4200

Epoch 15/30

Epoch 00015: sparse\_categorical\_accuracy improved from 0.44747 to 0.45484,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6366 - sparse\_categorical\_accuracy: 0.4548 - val\_loss:  
1.7256 - val\_sparse\_categorical\_accuracy: 0.4204

Epoch 16/30

Epoch 00016: sparse\_categorical\_accuracy improved from 0.45484 to 0.45648,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6256 - sparse\_categorical\_accuracy: 0.4565 - val\_loss:  
1.7474 - val\_sparse\_categorical\_accuracy: 0.4183

Epoch 17/30

Epoch 00017: sparse\_categorical\_accuracy improved from 0.45648 to 0.46274,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6134 - sparse\_categorical\_accuracy: 0.4627 - val\_loss:  
1.7137 - val\_sparse\_categorical\_accuracy: 0.4282

Epoch 18/30

Epoch 00018: sparse\_categorical\_accuracy improved from 0.46274 to 0.46619,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.6060 - sparse\_categorical\_accuracy: 0.4662 - val\_loss:  
1.7100 - val\_sparse\_categorical\_accuracy: 0.4299

Epoch 19/30

Epoch 00019: sparse\_categorical\_accuracy improved from 0.46619 to 0.46838,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5976 - sparse\_categorical\_accuracy: 0.4684 - val\_loss:  
1.7163 - val\_sparse\_categorical\_accuracy: 0.4245

Epoch 20/30

Epoch 00020: sparse\_categorical\_accuracy improved from 0.46838 to 0.46939,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5894 - sparse\_categorical\_accuracy: 0.4694 - val\_loss:  
1.7073 - val\_sparse\_categorical\_accuracy: 0.4300

Epoch 21/30

Epoch 00021: sparse\_categorical\_accuracy improved from 0.46939 to 0.47344,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5813 - sparse\_categorical\_accuracy: 0.4734 - val\_loss:  
1.7173 - val\_sparse\_categorical\_accuracy: 0.4271

Epoch 22/30

Epoch 00022: sparse\_categorical\_accuracy improved from 0.47344 to 0.47606,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5746 - sparse\_categorical\_accuracy: 0.4761 - val\_loss:  
1.7125 - val\_sparse\_categorical\_accuracy: 0.4272

Epoch 23/30

Epoch 00023: sparse\_categorical\_accuracy improved from 0.47606 to 0.47948,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5713 - sparse\_categorical\_accuracy: 0.4795 - val\_loss:  
1.7145 - val\_sparse\_categorical\_accuracy: 0.4255

Epoch 24/30

Epoch 00024: sparse\_categorical\_accuracy did not improve from 0.47948  
1946/1946 - 2s - loss: 1.5655 - sparse\_categorical\_accuracy: 0.4782 - val\_loss:  
1.6908 - val\_sparse\_categorical\_accuracy: 0.4336

Epoch 25/30

Epoch 00025: sparse\_categorical\_accuracy improved from 0.47948 to 0.48280,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5572 - sparse\_categorical\_accuracy: 0.4828 - val\_loss:  
1.7005 - val\_sparse\_categorical\_accuracy: 0.4343

Epoch 26/30

Epoch 00026: sparse\_categorical\_accuracy improved from 0.48280 to 0.48490,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5539 - sparse\_categorical\_accuracy: 0.4849 - val\_loss:  
1.7110 - val\_sparse\_categorical\_accuracy: 0.4290

Epoch 27/30

Epoch 00027: sparse\_categorical\_accuracy improved from 0.48490 to 0.48662,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5481 - sparse\_categorical\_accuracy: 0.4866 - val\_loss:  
1.6906 - val\_sparse\_categorical\_accuracy: 0.4344  
Epoch 28/30

Epoch 00028: sparse\_categorical\_accuracy did not improve from 0.48662  
1946/1946 - 2s - loss: 1.5445 - sparse\_categorical\_accuracy: 0.4857 - val\_loss:  
1.7195 - val\_sparse\_categorical\_accuracy: 0.4278  
Epoch 29/30

Epoch 00029: sparse\_categorical\_accuracy improved from 0.48662 to 0.48934,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5387 - sparse\_categorical\_accuracy: 0.4893 - val\_loss:  
1.6810 - val\_sparse\_categorical\_accuracy: 0.4387  
Epoch 30/30

Epoch 00030: sparse\_categorical\_accuracy improved from 0.48934 to 0.49035,  
saving model to ./checkpoints\_mlp\_best\_only/checkpoint  
1946/1946 - 2s - loss: 1.5345 - sparse\_categorical\_accuracy: 0.4903 - val\_loss:  
1.6972 - val\_sparse\_categorical\_accuracy: 0.4345

```
[40]: fig = plt.figure(figsize=(12, 5))

fig.add_subplot(121)

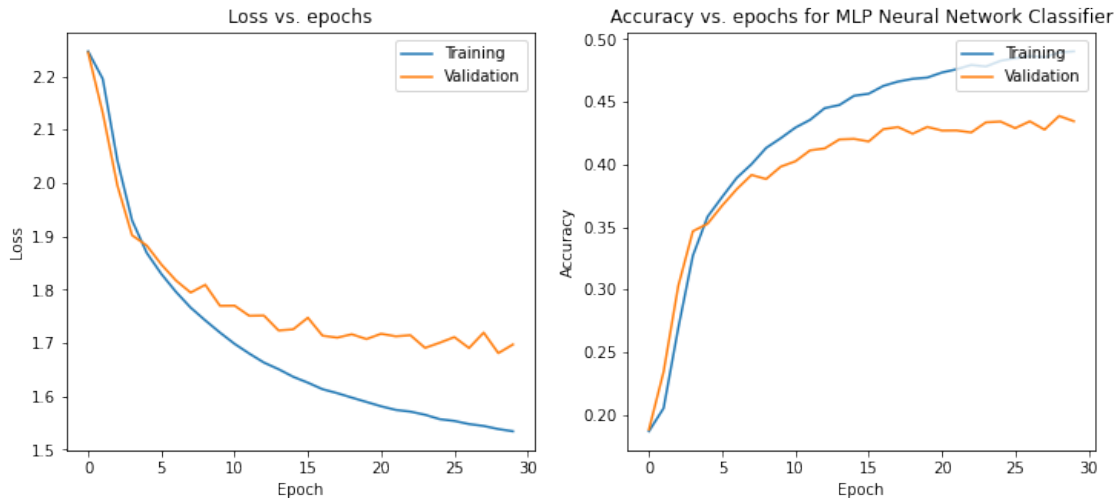
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

fig.add_subplot(122)

plt.plot(history.history['sparse_categorical_accuracy'])
plt.plot(history.history['val_sparse_categorical_accuracy'])
plt.title('Accuracy vs. epochs for MLP Neural Network Classifier')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()
```





```
[41]: # Compute and display the loss and accuracy of the trained model on the test_
      ↪ set.
test_loss, test_accuracy = model.evaluate(test_img_arr_gray, test_img_labels)
print(f"Test loss: {test_loss}")
print(f"Test accuracy: {test_accuracy}")
```

```
814/814 [=====] - 0s 552us/step - loss: 1.7390 -
sparse_categorical_accuracy: 0.4195
Test loss: 1.739003300666809
Test accuracy: 0.41952213644981384
```

### 1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.)*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[42]: from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,
      ↪BatchNormalization, Dropout
      from tensorflow.keras import regularizers

[47]: def create_model_cnn(input_shape, wd, rate):
      model = Sequential([
          Conv2D(32, (3,3), activation='relu', padding='SAME',
      ↪input_shape=input_shape),
          MaxPooling2D((2,2)),
          Dense(32, kernel_regularizer=regularizers.l2(wd), activation='relu'),
          BatchNormalization(),
          Dropout(rate),
          Dense(16, kernel_regularizer=regularizers.l2(wd), activation='relu'),
          Flatten(),
          Dense(11, activation='softmax', name='output_layer')])
      return model
      # create the CNN model
      model_cnn = create_model_cnn(input_shape, 1e-5, 0.3)
      # Print the model summary
      model_cnn.summary()
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 32, 32, 32)	320
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 32)	0
dense_4 (Dense)	(None, 16, 16, 32)	1056
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 32)	128
dropout_2 (Dropout)	(None, 16, 16, 32)	0
dense_5 (Dense)	(None, 16, 16, 16)	528
flatten_5 (Flatten)	(None, 4096)	0
output_layer (Dense)	(None, 11)	45067

Total params: 47,099  
 Trainable params: 47,035  
 Non-trainable params: 64

```
[48]: acc = tf.keras.metrics.SparseCategoricalAccuracy()
model_cnn.compile(optimizer="Adam", loss='sparse_categorical_crossentropy',
↳metrics=[acc])
# callbacks
checkpoint_best_only_cnn = get_checkpoint_best_only('./
↳checkpoints_cnn_best_only/checkpoint')
early_stopping_cnn = get_early_stopping()
callbacks_cnn = [checkpoint_best_only_cnn, early_stopping_cnn]
```

```
[49]: # Fit the model
history_acc = model_cnn.fit(train_img_arr_gray,
                             train_img_labels,
                             validation_split=0.15,
                             epochs=30,
                             batch_size=64,
                             verbose=2,
                             callbacks=callbacks_cnn)
```

Epoch 1/30

Epoch 00001: sparse\_categorical\_accuracy improved from -inf to 0.67643, saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 25s - loss: 1.0022 - sparse\_categorical\_accuracy: 0.6764 - val\_loss: 0.7903 - val\_sparse\_categorical\_accuracy: 0.7559

Epoch 2/30

Epoch 00002: sparse\_categorical\_accuracy improved from 0.67643 to 0.79262, saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.6800 - sparse\_categorical\_accuracy: 0.7926 - val\_loss: 0.7206 - val\_sparse\_categorical\_accuracy: 0.7826

Epoch 3/30

Epoch 00003: sparse\_categorical\_accuracy improved from 0.79262 to 0.80647, saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 25s - loss: 0.6373 - sparse\_categorical\_accuracy: 0.8065 - val\_loss: 0.7093 - val\_sparse\_categorical\_accuracy: 0.7847

Epoch 4/30

Epoch 00004: sparse\_categorical\_accuracy improved from 0.80647 to 0.81461, saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.6081 - sparse\_categorical\_accuracy: 0.8146 - val\_loss: 0.6431 - val\_sparse\_categorical\_accuracy: 0.8096

Epoch 5/30

Epoch 00005: sparse\_categorical\_accuracy improved from 0.81461 to 0.82285, saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.5855 - sparse\_categorical\_accuracy: 0.8228 - val\_loss: 0.6709 - val\_sparse\_categorical\_accuracy: 0.7987

Epoch 6/30

Epoch 00006: sparse\_categorical\_accuracy improved from 0.82285 to 0.82569,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.5714 - sparse\_categorical\_accuracy: 0.8257 - val\_loss:  
0.6596 - val\_sparse\_categorical\_accuracy: 0.8033

Epoch 7/30

Epoch 00007: sparse\_categorical\_accuracy improved from 0.82569 to 0.82986,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.5552 - sparse\_categorical\_accuracy: 0.8299 - val\_loss:  
0.6572 - val\_sparse\_categorical\_accuracy: 0.8039

Epoch 8/30

Epoch 00008: sparse\_categorical\_accuracy improved from 0.82986 to 0.83613,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.5386 - sparse\_categorical\_accuracy: 0.8361 - val\_loss:  
0.7037 - val\_sparse\_categorical\_accuracy: 0.7914

Epoch 9/30

Epoch 00009: sparse\_categorical\_accuracy improved from 0.83613 to 0.83814,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 25s - loss: 0.5294 - sparse\_categorical\_accuracy: 0.8381 - val\_loss:  
0.6684 - val\_sparse\_categorical\_accuracy: 0.8036

Epoch 10/30

Epoch 00010: sparse\_categorical\_accuracy improved from 0.83814 to 0.84234,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.5181 - sparse\_categorical\_accuracy: 0.8423 - val\_loss:  
0.6539 - val\_sparse\_categorical\_accuracy: 0.8049

Epoch 11/30

Epoch 00011: sparse\_categorical\_accuracy improved from 0.84234 to 0.84323,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.5140 - sparse\_categorical\_accuracy: 0.8432 - val\_loss:  
0.6745 - val\_sparse\_categorical\_accuracy: 0.8007

Epoch 12/30

Epoch 00012: sparse\_categorical\_accuracy improved from 0.84323 to 0.84427,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.5035 - sparse\_categorical\_accuracy: 0.8443 - val\_loss:  
0.6556 - val\_sparse\_categorical\_accuracy: 0.8047

Epoch 13/30

Epoch 00013: sparse\_categorical\_accuracy improved from 0.84427 to 0.84816,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.4961 - sparse\_categorical\_accuracy: 0.8482 - val\_loss:  
0.6445 - val\_sparse\_categorical\_accuracy: 0.8090

Epoch 14/30

Epoch 00014: sparse\_categorical\_accuracy improved from 0.84816 to 0.84910,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 28s - loss: 0.4897 - sparse\_categorical\_accuracy: 0.8491 - val\_loss:  
0.6697 - val\_sparse\_categorical\_accuracy: 0.8084

Epoch 15/30

Epoch 00015: sparse\_categorical\_accuracy improved from 0.84910 to 0.84992,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.4845 - sparse\_categorical\_accuracy: 0.8499 - val\_loss:  
0.6595 - val\_sparse\_categorical\_accuracy: 0.8078

Epoch 16/30

Epoch 00016: sparse\_categorical\_accuracy improved from 0.84992 to 0.85196,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 26s - loss: 0.4793 - sparse\_categorical\_accuracy: 0.8520 - val\_loss:  
0.6737 - val\_sparse\_categorical\_accuracy: 0.7977

Epoch 17/30

Epoch 00017: sparse\_categorical\_accuracy improved from 0.85196 to 0.85333,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.4736 - sparse\_categorical\_accuracy: 0.8533 - val\_loss:  
0.6670 - val\_sparse\_categorical\_accuracy: 0.8041

Epoch 18/30

Epoch 00018: sparse\_categorical\_accuracy did not improve from 0.85333  
973/973 - 27s - loss: 0.4683 - sparse\_categorical\_accuracy: 0.8532 - val\_loss:  
0.6705 - val\_sparse\_categorical\_accuracy: 0.8033

Epoch 19/30

Epoch 00019: sparse\_categorical\_accuracy improved from 0.85333 to 0.85635,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.4648 - sparse\_categorical\_accuracy: 0.8563 - val\_loss:  
0.6967 - val\_sparse\_categorical\_accuracy: 0.7936

Epoch 20/30

Epoch 00020: sparse\_categorical\_accuracy improved from 0.85635 to 0.85813,  
saving model to ./checkpoints\_cnn\_best\_only/checkpoint  
973/973 - 27s - loss: 0.4575 - sparse\_categorical\_accuracy: 0.8581 - val\_loss:  
0.7199 - val\_sparse\_categorical\_accuracy: 0.7942

Epoch 21/30

Epoch 00021: sparse\_categorical\_accuracy did not improve from 0.85813  
973/973 - 27s - loss: 0.4569 - sparse\_categorical\_accuracy: 0.8573 - val\_loss:  
0.7166 - val\_sparse\_categorical\_accuracy: 0.7872

Epoch 22/30

Epoch 00022: sparse\_categorical\_accuracy did not improve from 0.85813  
 973/973 - 27s - loss: 0.4563 - sparse\_categorical\_accuracy: 0.8571 - val\_loss:  
 0.6871 - val\_sparse\_categorical\_accuracy: 0.7983

```
[50]: fig = plt.figure(figsize=(12, 5))

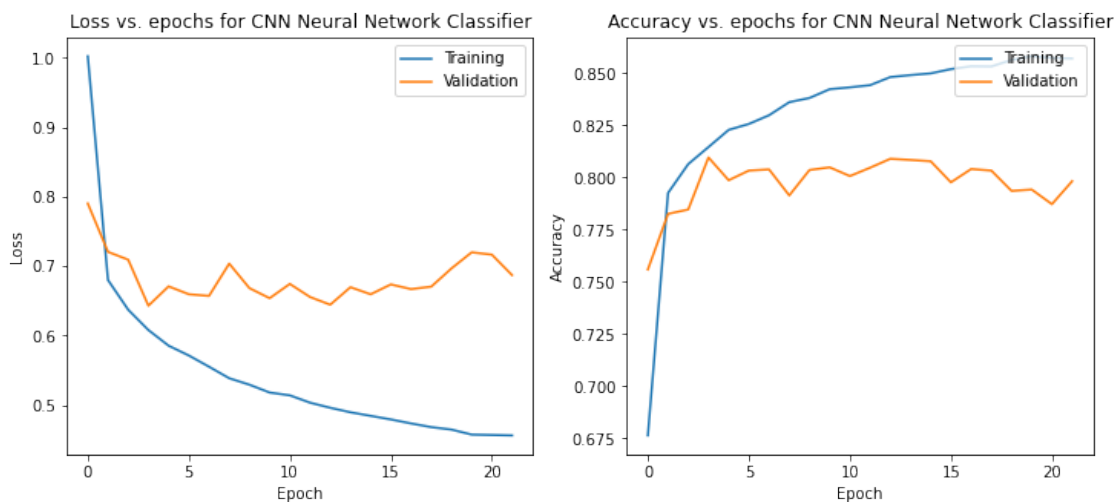
fig.add_subplot(121)

plt.plot(history_acc.history['loss'])
plt.plot(history_acc.history['val_loss'])
plt.title('Loss vs. epochs for CNN Neural Network Classifier')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

fig.add_subplot(122)

plt.plot(history_acc.history['sparse_categorical_accuracy'])
plt.plot(history_acc.history['val_sparse_categorical_accuracy'])
plt.title('Accuracy vs. epochs for CNN Neural Network Classifier')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()
```



```
[51]: # Compute and display the loss and accuracy of the trained model on the test_
      ↪ set.
test_loss, test_accuracy = model_cnn.evaluate(test_img_arr_gray,
      ↪ test_img_labels)
```

```
print(f"Test loss: {test_loss}")
print(f"Test accuracy: {test_accuracy}")
```

```
814/814 [=====] - 3s 4ms/step - loss: 0.7880 -
sparse_categorical_accuracy: 0.7708
Test loss: 0.788002610206604
Test accuracy: 0.7708205580711365
```

## 1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```
[52]: # get best weights for the MLP model
model.load_weights('checkpoints_mlp_best_only/checkpoint')
# get the beset weights for the CNN model
model_cnn.load_weights('checkpoints_cnn_best_only/checkpoint')
```

```
[52]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x16282ae50>
```

```
[53]: # get MLP and CNN model predictions on randomly selected test images

num_test_images = test_img_arr_gray.shape[0]

random_inx = np.random.choice(num_test_images, 5)
random_test_images_gray = test_img_arr_gray[random_inx, ...]
random_test_labels = test_img_labels[random_inx, ...]

predictions_mlp = model.predict(random_test_images_gray)
predictions_cnn = model_cnn.predict(random_test_images_gray)
```

```
[80]: # MLP predictions
fig, axes = plt.subplots(5, 2, figsize=(16, 12))

fig.subplots_adjust(hspace=0.4, wspace=-0.2)

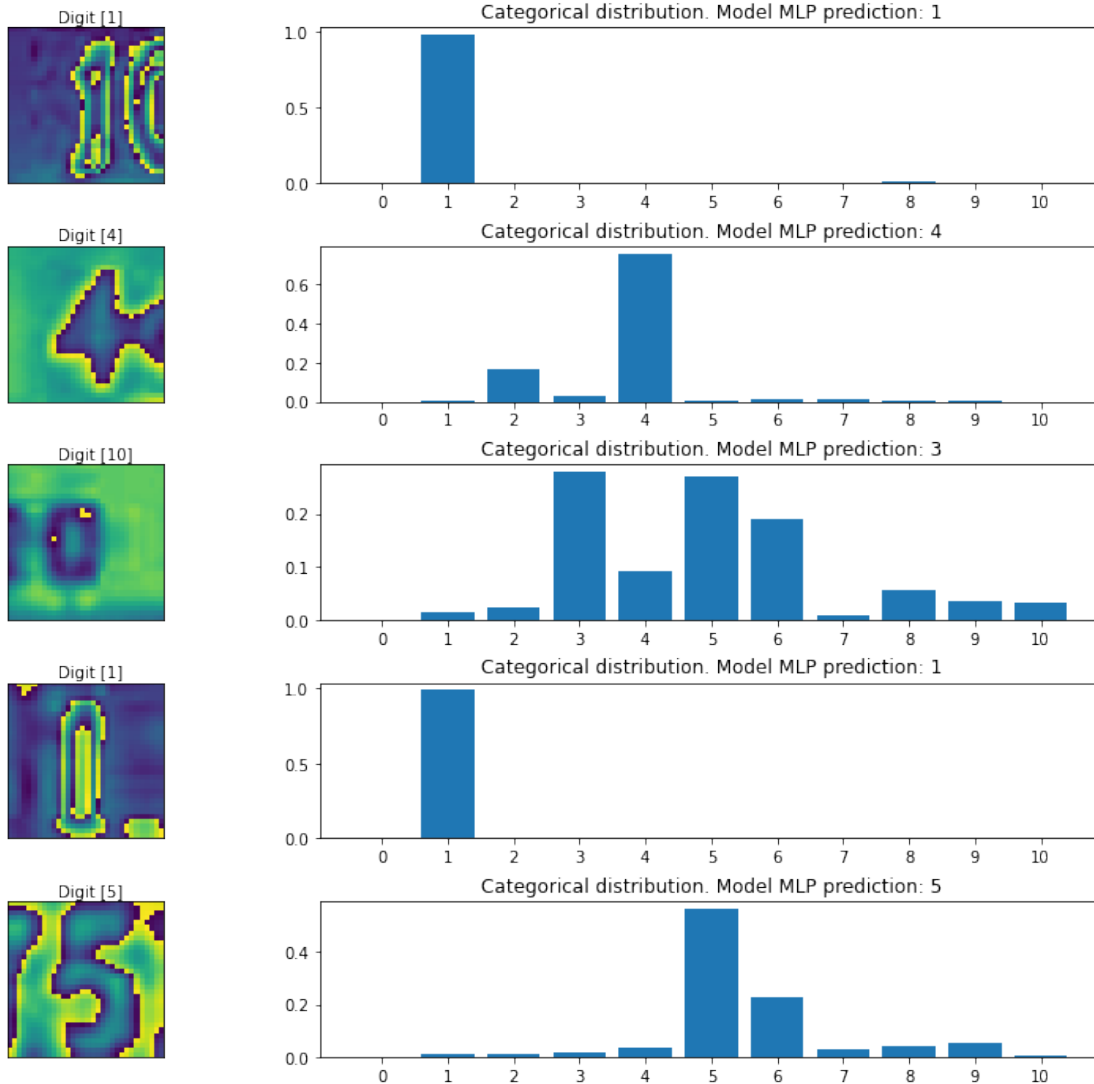
for i, (prediction_mlp, image, label) in enumerate(zip(predictions_mlp,
↳ random_test_images_gray, random_test_labels)):
    axes[i, 0].imshow(np.squeeze(image))
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10., -1.5, f'Digit {label}')
    axes[i, 1].bar(np.arange(len(prediction_mlp)), prediction_mlp)
    axes[i, 1].set_xticks(np.arange(len(prediction_mlp)))
```

```

    axes[i, 1].set_title(f"Categorical distribution. Model MLP prediction: {np.
↪argmax(prediction_mlp)}")

plt.show()

```



```

[55]: # CNN predictions
fig, axes = plt.subplots(5, 2, figsize=(16, 12))

fig.subplots_adjust(hspace=0.4, wspace=-0.2)

for i, (prediction_cnn, image, label) in enumerate(zip(predictions_cnn,
↪random_test_images_gray, random_test_labels)):
    axes[i, 0].imshow(np.squeeze(image))

```



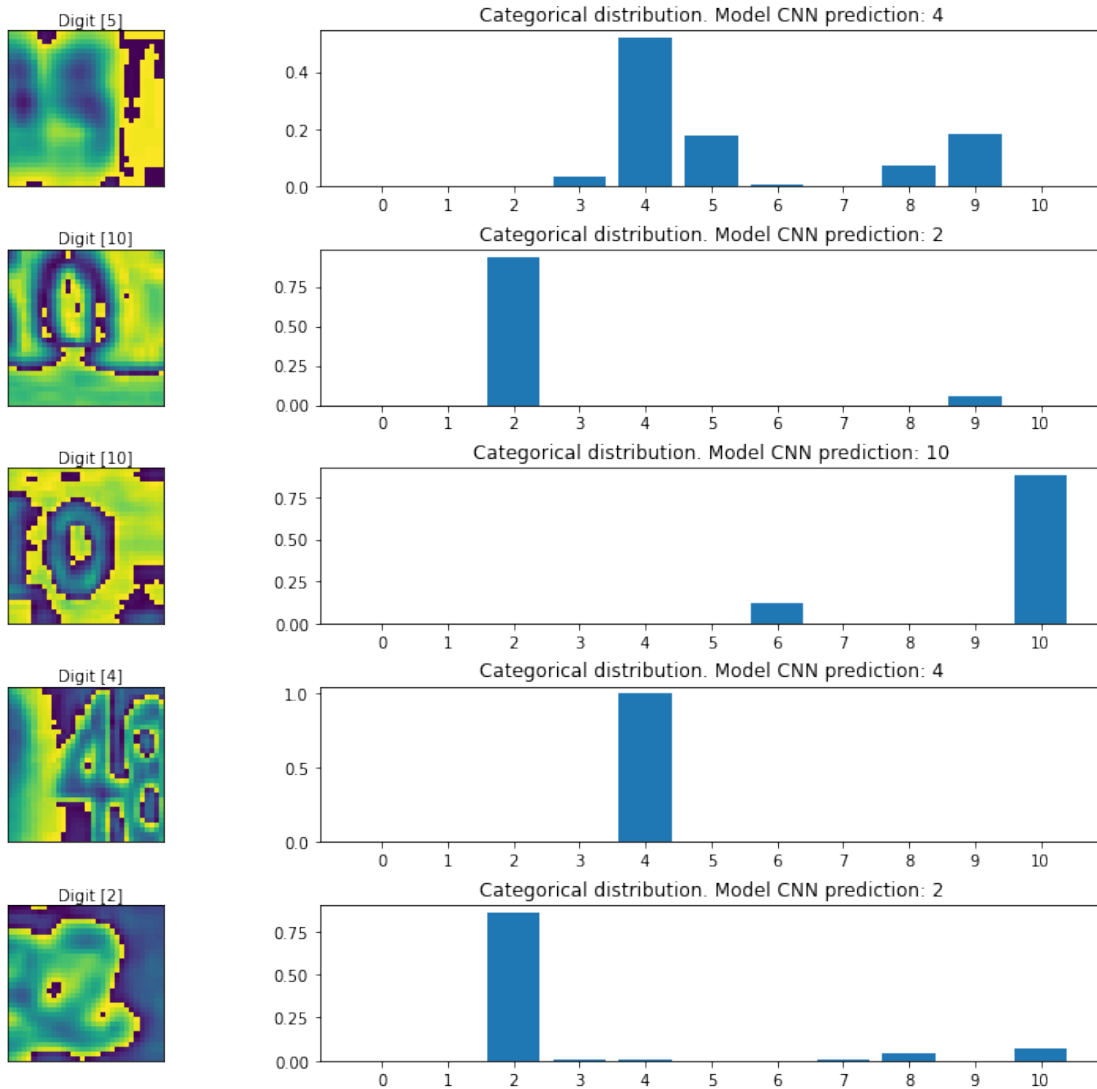
```

axes[i, 0].get_xaxis().set_visible(False)
axes[i, 0].get_yaxis().set_visible(False)
axes[i, 0].text(10., -1.5, f'Digit {label}')
```

```

axes[i, 1].bar(np.arange(len(prediction_cnn)), prediction_cnn)
axes[i, 1].set_xticks(np.arange(len(prediction_cnn)))
axes[i, 1].set_title(f"Categorical distribution. Model CNN prediction: {np.
↪argmax(prediction_cnn)}")

plt.show()
```



[82] :

[ ] :